

Data Mining (W4240 Section 001)

Boosting

Giovanni Motta

Columbia University, Department of Statistics

November 25, 2015

Outline

Ensemble Methods

Boosting

AdaBoost

Examples

Statistics of Boosting

AdaBoost and Logistic Regression

Summary

Outline

Ensemble Methods

Boosting

AdaBoost

Examples

Statistics of Boosting

AdaBoost and Logistic Regression

Summary

Model Fitting

What we have done:

- ▶ get data
- ▶ fit some models
- ▶ evaluate results
- ▶ choose best model, apply to test data

Last time and today we consider an alternative approach:

- ▶ get data
- ▶ fit some models
- ▶ evaluate results
- ▶ combine models, apply to test data

Approaches of this type are generally called *ensemble methods*

Ensemble Methods

Ensembles methods use collections of models to get better predictive performance than any single model

- ▶ get a collection of predictive models
- ▶ the ensemble predictor is an average of the underlying models
- ▶ we introduced bagging
- ▶ we introduced random forests

Why should this work?

- ▶ often easy to fit simple models well
- ▶ if we average lots of *different* simple models, we can fit these well and have a large model space
- ▶ and we can reduce the variance of the estimator

And the gun fires (foreshadowing from last time)

- ▶ Bagging: we average all our estimators together
- ▶ Don't we believe that some will be more useful than others?
- ▶ Could we create a weighted average of estimators?
- ▶ Can we learn better estimators as we go?
- ▶ This time...

Outline

Ensemble Methods

Boosting

AdaBoost

Examples

Statistics of Boosting

AdaBoost and Logistic Regression

Summary

Strong Learners vs. Weak Learners

Setting: classification with 2 classes (here, -1 and 1)

A **strong learner** is a method that can learn a decision rule arbitrarily well.

A **weak learner** is a simple method that does better than guessing, but cannot learn a decision rule arbitrarily well.

Example: trying to decide whether an email is ham or spam.

- ▶ strong learner example: method that uses words, syntax, etc as features, and fits a high-accuracy decision rule
- ▶ weak learner example: “If the phrase ‘lose weight’ is in the email, then predict it is spam”

Boosting

Can we combine weak learners to make a strong learner? If so, how can we do it in a computationally efficient manner?

Boosting

Can we combine weak learners to make a strong learner? If so, how can we do it in a computationally efficient manner?

Answer: yes, we can combine weak learners to make a strong learner with *boosting*

Boosting:

- ▶ start with a method for finding weak learners
- ▶ call this method repeatedly, each time with *new* subsets of the data
- ▶ The i th subset is a random sample of the data with some weight (ex: $p_1 = \frac{1}{2n}$, $p_2 = \frac{1}{1.5n}$, ...)
- ▶ new predictor is a weighted average of the weak learners
- ▶ (note the implicit idea that we are educating ourselves as we go, and weighting more as we learn more)

Boosting

Assuming we have a good way to generate weak learners (success rate more than 50%), we still have some problems:

- ▶ how should we reweight the data each round?
- ▶ how should we combine the weak rules into a single (strong?) rule?

Want:

- ▶ want higher weights on data that have been previously misclassified
- ▶ prediction to be a simple weighted majority of rules

Outline

Ensemble Methods

Boosting

AdaBoost

Examples

Statistics of Boosting

AdaBoost and Logistic Regression

Summary

AdaBoost

AdaBoost (**Adaptive Boosting**) is one way to do this.

Data: $(x_i, y_i)_{i=1}^n$, x_i is a vector and $y_i \in \{-1, 1\}$

Inputs: data, number of rounds T , weak learner $\hat{y} = h_t(x)$

Output: decision rule H

Weights:

▶ start with $\mathcal{D}_1(i) = \frac{1}{n}$ for $i = 1, \dots, n$

▶ for $t = 1, \dots, T$:

$=1 \parallel =-1$ (i.e. prod=1)

$$\begin{aligned}\mathcal{D}_{t+1}(i) &= \frac{\mathcal{D}_t(i)}{Z_t} \times \begin{cases} e^{-\alpha_t} & \text{if } y_i = h_t(x_i) \text{ (smaller weights for easy examples)} \\ e^{\alpha_t} & \text{if } y_i \neq h_t(x_i) \text{ (larger weights for hard examples)} \end{cases} \\ &= \frac{\mathcal{D}_t(i)}{Z_t} e^{-y_i \alpha_t h_t(x_i)}\end{aligned}$$

(here Z_t is a normalization constant so weights sum to 1)
(who can tell me why the second equality holds?)

AdaBoost

Weights: this is an exponential weighting scheme. Easy examples are downweighted, hard examples are upweighted.

What about the predictor, H ? It is a weighted linear combination of the weak learners,

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$

The weights assigned are directly related to how well h_t performed on the weighted training set:

$$\alpha_t = \frac{1}{2} \log \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$$

where

$$\epsilon_t = \mathbb{P}[h_t(x_i) \neq y_i] \approx \sum_{i=1}^n \mathcal{D}_t(i) \mathbf{1}_{\{h_t(x_i) \neq y_i\}}$$

AdaBoost: Forward stagewise additive modeling

Consider the problem of fitting a basis function of the form

Final Goal is $\text{sgn}(f)$
Use Fourier Transform
Update everyday!
But $T=30\sim 40$

$$f(x) = \sum_{t=1}^T \alpha_t h_t(x)$$

- ▶ data set $\{(x_1, y_1), \dots, (x_n, y_n)\}$ where $y_i \in \{-1, 1\}$
- ▶ set of weak classifiers $\{h_1, \dots, h_T\}$ each of which outputs a classification $h_t(x_i) \in \{-1, 1\}$ for each item
- ▶ At iteration $t - 1$, our boosted classifier is the solution of minimize loss function

$$(\alpha_t, h_t) = \arg \min_{\alpha, h} \sum_{i=1}^n L[y_i, f_{(t-1)}(x_i) + \alpha h(x_i)]$$

- ▶ At the t -th iteration, we want to extend this to a better boosted classifier:

$$f_t(x_i) = f_{(t-1)}(x_i) + \alpha_t h_t(x_i)$$

AdaBoost: Forward stagewise additive modeling

Choose the exponential loss function: $L[y, f(x)] = \exp[-y f(x)]$

- At iteration $t - 1$, our boosted classifier is the solution of

$$\begin{aligned}(\alpha_t, h_t) &= \arg \min_{\alpha, h} \sum_{i=1}^n \exp \{ -y_i [f_{(t-1)}(x_i) + \alpha h(x_i)] \} \\ &= \arg \min_{\alpha, h} \sum_{i=1}^n \mathcal{D}_t(i) \exp \{ -\alpha y_i h(x_i) \} \quad (1)\end{aligned}$$

where $\mathcal{D}_t(i) = \exp[-y_i f_{(t-1)}(x_i)]$.

- At the t -th iteration

$$f_t(x_i) = f_{(t-1)}(x_i) + \alpha_t h_t(x_i)$$

The solution to (1) is

$$\begin{aligned}h_t &= \arg \min_h \sum_{i=1}^n \mathcal{D}_t(i) \mathbf{1}_{y_i \neq h(x_i)} \\ \alpha_t &= \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)\end{aligned}$$

AdaBoost

Algorithm:

- ▶ Initialize weights $\mathcal{D}_1(i) = \frac{1}{n}$ for $i = 1, \dots, n$
- ▶ For each round $t = 1, \dots, T$:
 - ▶ draw a “sufficiently large” sample i.i.d. from \mathcal{D}_t
 - ▶ run weak learning algorithm on this sample to produce rule h_t
 - ▶ set $\epsilon_t = \sum_{i=1}^n \mathcal{D}_t(i) \mathbf{1}_{\{h_t(x_i) \neq y_i\}}$
 - ▶ set $\alpha_t = \frac{1}{2} \log \left(\frac{1-\epsilon_t}{\epsilon_t} \right)$
 - ▶ update distribution by setting

$$\mathcal{D}_{t+1}(i) = \frac{\mathcal{D}_t(i)}{Z_t} e^{-\alpha_t y_i h_t(x_i)}$$

- ▶ Output function H where

$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$

AdaBoost

Why AdaBoost?

- ▶ no tunable parameters Unlike Ridge and Lasso has λ
- ▶ don't need to know how well weak learners do
- ▶ works with any weak learner
- ▶ computationally reasonable
- ▶ tends to avoid overfitting
- ▶ has some nice theoretical guarantees

Pause: **A hugely important method, running all over the world**

Getting Learners

So how can we get weak learners?

- ▶ use weak learners that are actually pretty strong—they come from another learning algorithm

- ▶ decision trees
- ▶ naive Bayes
- ▶ k-nn
- ▶ generalized linear models
- ▶ splines
- ▶ etc

Example: decision trees

Each round, AdaBoost gives the algorithm a new set of data and asks it to come up with a new tree with low error. Each tree is an h_t .

Outline

Ensemble Methods

Boosting

AdaBoost

Examples

Statistics of Boosting

AdaBoost and Logistic Regression


Summary

A simple example

- ▶ 10d data: $X_1, \dots, X_{10} \sim_{iid} \mathcal{N}(0, 1)$

$$Y = \begin{cases} +1 & \text{if } \sum_{i=1}^{10} X_i^2 > \chi_{10}^2(0.5) = 9.34, \\ -1 & \text{otherwise} \end{cases}$$


- ▶ Classifier is a one-level tree (stump)
- ▶ 2000 training points (roughly 50/50), and 10,000 test points
- ▶ Error by boosting iterations:



- ▶ (let's interpret this classifier...)

A simple example

- ▶ Red: misclassification error rate on the training set (red)
- ▶ Blue: average exponential loss $\frac{1}{n} \sum_{i=1}^n \exp[-y_i f(x_i)]$.



Boosting in R

There are many, many packages that implement boosting in R

- ▶ `ada`: stumpy trees from CART
- ▶ `gmb`: boosting for regression with trees
- ▶ `mboost`: model-based boosting for regression and classification with a variety of methods, including GLMs and splines
- ▶ many more....

Boosting in R

Let's use the package `ada` on the `kyphosis` dataset (from `rpart`).

```
> library(rpart)
> library(ada)
> dim(kyphosis)
> ind.train <- sample(1:81,60)
> ind.test <- setdiff(1:81,ind.train)
> # iter is number of iterations
> # nu is a shrinkage parameter
> ky.ada <- ada(Kyphosis~.,data=kyphosis[ind.train,],iter=20,nu=1,type="discrete")
> # add testing data set
> ky.ada.test <- addtest(ky.ada,kyphosis[ind.test,-1],kyphosis[ind.test,1])
> plot(ky.ada.test,TRUE,TRUE)
> ky.ada.test2 <- predict(ky.ada,kyphosis[ind.test,-1],type="vector")
> # Make a tree
> ky.rpart <- rpart(Kyphosis~.,data=kyphosis[ind.train,])
> ky.rpart.test <- predict(ky.rpart,kyphosis[ind.test,-1],type="class")
> cbind(kyphosis[ind.test,1],ky.rpart.test,ky.ada.test2)
```


Outline

Ensemble Methods

Boosting

AdaBoost

Examples


Statistics of Boosting

AdaBoost and Logistic Regression

Summary

A Statistical View of Boosting

Simplified version: AdaBoost is a coordinate descent algorithm



Coordinate descent:

- ▶ have function f ; take gradient
- ▶ move in the *coordinate direction* with largest change
- ▶ length of move is α_t ; repeat

A Statistical View of Boosting

Misclassification error for some function f :

$$\frac{1}{n} \sum_{i=1}^n \mathbf{1}_{\{y_i f(x_i) \leq 0\}}$$

The misclassification error is upper bounded by the exponential loss:

$$\frac{1}{n} \sum_{i=1}^n e^{-y_i f(x_i)}$$

...and we want f to be a linear combination of classifiers,

$$f(x) = \sum_{j=1}^m \lambda_j h_j(x)$$

So we will minimize the exponential loss with respect to the λ_j 's.

A Statistical View of Boosting

Here are the details:

- ▶ define an $n \times m$ matrix M with $M_{ij} = y_i h_j(x_i) \in \{\pm 1\}$

$$M = \begin{matrix} & \begin{matrix} \text{weak classifiers} \\ j \end{matrix} \\ \begin{matrix} \text{examples} \\ i \end{matrix} & \begin{bmatrix} \pm 1 \end{bmatrix} \end{matrix}$$

- ▶ then

$$y_i f(x_i) = \sum_j \lambda_j y_i h_j(x_i) = \sum_j \lambda_j M_{ij} = (M\lambda)_i$$

- ▶ which has the exponential loss

$$R^{train}(\lambda) = \frac{1}{n} \sum_i e^{-y_i f(x_i)} = \frac{1}{n} \sum_i e^{-(M\lambda)_i}$$

A Statistical View of Boosting

Let's do coordinate descent on the exponential loss, R^{train} .

- ▶ each iteration, choose a coordinate of λ , called j_t , and move α_t in the j^{th} direction (classifier gives direction, α_t length of step)
- ▶ find direction with steepest gradient: (set \mathbf{e}_j as vector of 0's with 1 in j^{th} place)

$$\begin{aligned}j_t &\in \arg \max_j \left[- \left. \frac{\partial R^{train}(\lambda_t + \alpha \mathbf{e}_j)}{\partial \alpha} \right|_{\alpha=0} \right] \\&= \arg \max_j \left[- \frac{\partial}{\partial \alpha} \left[\frac{1}{n} \sum_{i=1}^n e^{-(M(\lambda_t + \alpha \mathbf{e}_j))_i} \right] \right|_{\alpha=0} \right] \\&= \arg \max_j \left[- \frac{\partial}{\partial \alpha} \left[\frac{1}{n} \sum_{i=1}^n e^{-(M\lambda_t)_i - \alpha(M\mathbf{e}_j)_i} \right] \right|_{\alpha=0} \right] \\&= \arg \max_j \left[- \frac{\partial}{\partial \alpha} \left[\frac{1}{n} \sum_{i=1}^n e^{-(M\lambda_t)_i - \alpha M_{ij}} \right] \right|_{\alpha=0} \right] \\&= \arg \max_j \left[\frac{1}{n} \sum_{i=1}^n M_{ij} e^{-(M\lambda_t)_i} \right]\end{aligned}$$

A Statistical View of Boosting

To deal with

$$j_t \in \arg \max_j \left[\frac{1}{n} \sum_{i=1}^n M_{ij} e^{-(M\lambda_t)_i} \right],$$

create a probability distribution

$$\mathcal{D}_t(i) = e^{-(M\lambda_t)_i} / Z_t, \text{ where } Z_t = \sum_{i=1}^n e^{-(M\lambda_t)_i}$$

Since we are just multiplying by constants,

$$j_t \in \arg \max_j \sum_{i=1}^n M_{ij} \mathcal{D}_t(i) = \arg \max_j (\mathcal{D}_t^T M)_j$$

So that's how we choose the classifier at step t .

A Statistical View of Boosting

So now that we have chosen j_t , how far should we move in that direction? (i.e., how much of that classifier should we add in to the large classifier?)

$$\begin{aligned} 0 &= \left. \frac{\partial R^{train}(\lambda_t + \alpha \mathbf{e}_{j_t})}{\partial \alpha} \right|_{\alpha_t} \\ &= -\frac{1}{n} \sum_{i=1}^n M_{ijt} e^{-(M\lambda_t)_i - \alpha_t M_{ijt}} \\ &= -\frac{1}{n} \sum_{i: M_{ijt}=1} e^{-(M\lambda_t)_i} e^{-\alpha_t} - \frac{1}{n} \sum_{i: M_{ijt}=-1} -e^{-(M\lambda_t)_i} e^{\alpha_t} \\ &= \sum_{i: M_{ijt}=1} \mathcal{D}_t(i) e^{-\alpha_t} - \sum_{i: M_{ijt}=-1} \mathcal{D}_t(i) e^{\alpha_t} \\ &=: d_+ e^{-\alpha_t} - d_- e^{\alpha_t} \\ e^{2\alpha_t} &= \frac{d_+}{d_-}, \quad \alpha_t = \frac{1}{2} \log \frac{d_+}{d_-} = \frac{1}{2} \log \frac{1 - d_-}{d_-} \end{aligned}$$

A Statistical View of Boosting

The new coordinate descent algorithm is:

- ▶ set $\mathcal{D}_1(i) = \frac{1}{n}$ for $i = 1, \dots, n$; $\lambda_1 = 0$
- ▶ for $t = 1, \dots, T$:
 - ▶ set $j_t \in \arg \max_j (\mathcal{D}_t^T M)_j$
 - ▶ set $d_- = \sum_{M_{ij_t} = -1} \mathcal{D}_t(i)$
 - ▶ set $\alpha_t = \frac{1}{2} \log \frac{1-d_-}{d_-}$
 - ▶ set $\lambda_{t+1} = \lambda_t + \alpha_t \mathbf{e}_{j_t}$
 - ▶ set $\mathcal{D}_{t+1}(i) = e^{-(M\lambda_{t+1})_i} / Z_t$ for each i
 - ▶ set $Z_t = \sum_{i=1}^n e^{-(M\lambda_{t+1})_i}$
- ▶ set $f(x) = \sum_{j=1}^m \lambda_j h_j(x)$

A Statistical View of Boosting

The new coordinate descent algorithm is:

- ▶ set $\mathcal{D}_1(i) = \frac{1}{n}$ for $i = 1, \dots, n$; $\lambda_1 = 0$
- ▶ for $t = 1, \dots, T$:
 - ▶ set $j_t \in \arg \max_j (\mathcal{D}_t^T M)_j$
 - ▶ set $d_- = \sum_{M_{ij_t} = -1} \mathcal{D}_t(i)$
 - ▶ set $\alpha_t = \frac{1}{2} \log \frac{1-d_-}{d_-}$
 - ▶ set $\lambda_{t+1} = \lambda_t + \alpha_t \mathbf{e}_{j_t}$
 - ▶ set $\mathcal{D}_{t+1}(i) = e^{-(M\lambda_{t+1})_i} / Z_t$ for each i
 - ▶ set $Z_t = \sum_{i=1}^n e^{-(M\lambda_{t+1})_i}$
- ▶ set $f(x) = \sum_{j=1}^m \lambda_j h_j(x)$

Is this AdaBoost?

A Statistical View of Boosting

How does λ_t relate to α_t ?

- ▶ $\lambda_{T,j}$ is the sum of the α_t 's where the chosen direction is j

$$\lambda_{T,j} = \sum_{t=1}^T \alpha_t \mathbf{1}_{\{j_t=j\}}$$

The output function f is

$$\begin{aligned} f(x) &= \sum_{j=1}^m \sum_{t=1}^T \lambda_{t,j} h_j(x) \\ &= \sum_{j=1}^m \sum_{t=1}^T \alpha_t \mathbf{1}_{\{j_t=j\}} h_j(x) \\ &= \sum_{t=1}^T \alpha_t \sum_{j=1}^m h_j(x) \mathbf{1}_{\{j_t=j\}} \\ &= \sum_{t=1}^T \alpha_t h_{j_t}(x) \end{aligned}$$

A Statistical View of Boosting

OK, so the output function has a similar form. What about the weights?

Weights for AdaBoost:

$$\begin{aligned}\mathcal{D}_{t+1}(i) &= \frac{\mathcal{D}_t(i)e^{-M_{ij_t}\alpha_t}}{Z_t} &&= \frac{\prod_t e^{-M_{ij_t}\alpha_t}}{n \prod_t Z_t} \\ &= \frac{e^{-\sum_t M_{ij_t}\alpha_t}}{n \prod_t Z_t} &&= \frac{1}{n \prod_t Z_t} e^{-\sum_j M_{ij}\lambda_{t,j}}\end{aligned}$$

The denominator must be $\sum_i e^{-\sum_j M_{ij}\lambda_{t,j}}$ since the weights sum to 1.

Therefore the weights are the same, as long as the j_t 's and α_t 's are the same.

A Statistical View of Boosting

Let's check the chosen directions, the j_t 's. AdaBoost chooses the classifier with the lowest error. In terms of the matrix M :

$$\begin{aligned}j_t &\in \arg \min_j \sum_i \mathcal{D}_t(i) \mathbf{1}_{\{h_j(x_i) \neq y_i\}} \\&= \arg \min_j \sum_{i: M_{ij} = -1} \mathcal{D}_t(i) \\&= \arg \max_j \left[- \sum_{i: M_{ij} = -1} \mathcal{D}_t(i) \right] \\&= \arg \max_j \left[1 - 2 \sum_{i: M_{ij} = -1} \mathcal{D}_t(i) \right] \\&= \arg \max_j \left[\left[\sum_{i: M_{ij} = 1} \mathcal{D}_t(i) + \sum_{i: M_{ij} = -1} \mathcal{D}_t(i) \right] - 2 \sum_{i: M_{ij} = -1} \mathcal{D}_t(i) \right] \\&= \arg \max_j \sum_{i: M_{ij} = 1} \mathcal{D}_t(i) - \sum_{i: M_{ij} = -1} \mathcal{D}_t(i) = \arg \max_j (\mathcal{D}_t^T M)_j\end{aligned}$$

A Statistical View of Boosting

OK, so we have the same weights and the same chosen direction.
Are the step sizes the same? AdaBoost error rate:

$$\begin{aligned}\epsilon_t &= \sum_i \mathcal{D}_t(i) \mathbf{1}_{\{h_{j_t}(x_i) \neq y_i\}} \\ &= \sum_{i: h_{j_t}(x_i) \neq y_i} \mathcal{D}_t(i) \\ &= \sum_{i: M_{ij_t} = -1} \mathcal{D}_t(i) = d_-\end{aligned}$$

Working our way from the AdaBoost stepsize:

$$\begin{aligned}\alpha_t &= \frac{1}{2} \log \frac{1 - \epsilon_t}{\epsilon_t} \\ &= \frac{1}{2} \log \frac{1 - d_-}{d_-}\end{aligned}$$

A Statistical View of Boosting

Conclusion: AdaBoost minimizes exponential loss using coordinate descent.

Outline

Ensemble Methods

Boosting

AdaBoost

Examples

Statistics of Boosting

AdaBoost and Logistic Regression

Summary

AdaBoost and Logistic Regression

Logistic regression has a logistic loss function. How is it similar to AdaBoost?

- ▶ AdaBoost approximately minimizes exponential loss over the whole distribution
- ▶ We can use this idea to get label probabilities from AdaBoost

Lemma (Hastie, Friedman, Tibshirani, 2001)

$$\mathbb{E}_{Y \sim \mathcal{D}(x)} e^{-Y f(x)}$$

is minimized at

$$f(x) = \frac{1}{2} \log \frac{\mathbb{P}(Y = 1 | x)}{\mathbb{P}(Y = -1 | x)};$$

AdaBoost and Logistic Regression

Proof.

$$\begin{aligned}\mathbb{E}e^{-Yf(x)} &= \mathbb{P}(Y = 1 | x)e^{-f(x)} + \mathbb{P}(Y = -1 | x)e^{f(x)} \\ 0 &= \frac{d \mathbb{E}(e^{-Yf(x)} | x)}{d f(x)} = -\mathbb{P}(Y = 1 | x)e^{-f(x)} + \mathbb{P}(Y = -1 | x)e^{f(x)} \\ \mathbb{P}(Y = 1 | x)e^{-f(x)} &= \mathbb{P}(Y = -1 | x)e^{f(x)} \\ \frac{\mathbb{P}(Y = 1 | x)}{\mathbb{P}(Y = -1 | x)} &= e^{2f(x)} \\ f(x) &= \frac{1}{2} \log \frac{\mathbb{P}(Y = 1 | x)}{\mathbb{P}(Y = -1 | x)}\end{aligned}$$



AdaBoost and Logistic Regression

How does this compare with logistic regression?

Logistic regression:

$$f(x) = \log \frac{\mathbb{P}(Y = 1 | x)}{\mathbb{P}(Y = -1 | x)}$$

so we are off by a factor of 2.

Use this fact and the previous lemma to get probabilities out of AdaBoost by solving for $p = \mathbb{P}(Y = 1 | x)$.

AdaBoost and Logistic Regression

$$f(x) = \frac{1}{2} \log \frac{\mathbb{P}(Y = 1 | x)}{\mathbb{P}(Y = -1 | x)} =: \frac{1}{2} \log \frac{p}{1-p}$$

$$e^{2f(x)} = \frac{p}{1-p}$$

$$e^{2f(x)} - pe^{2f(x)} = p$$

$$e^{2f(x)} = p(1 + e^{2f(x)})$$

$$p = \mathbb{P}(Y = 1 | x) = \frac{e^{2f(x)}}{1 + e^{2f(x)}}$$

Even though AdaBoost minimizes a different objective function, the probability of success is similar to that of logistic regression (remove the 2's).

Outline

Ensemble Methods

Boosting

AdaBoost

Examples

Statistics of Boosting

AdaBoost and Logistic Regression

Summary

Boosting

What I want you to get from this:

- ▶ connection between weak learners and strong learners
- ▶ what the algorithm is (you will be implementing it in HW)
- ▶ general idea of why it works