




Cluster Analysis II

MENGQIAN LU

Assess your clusters with visualization

`silhouette()` in `library(cluster)`

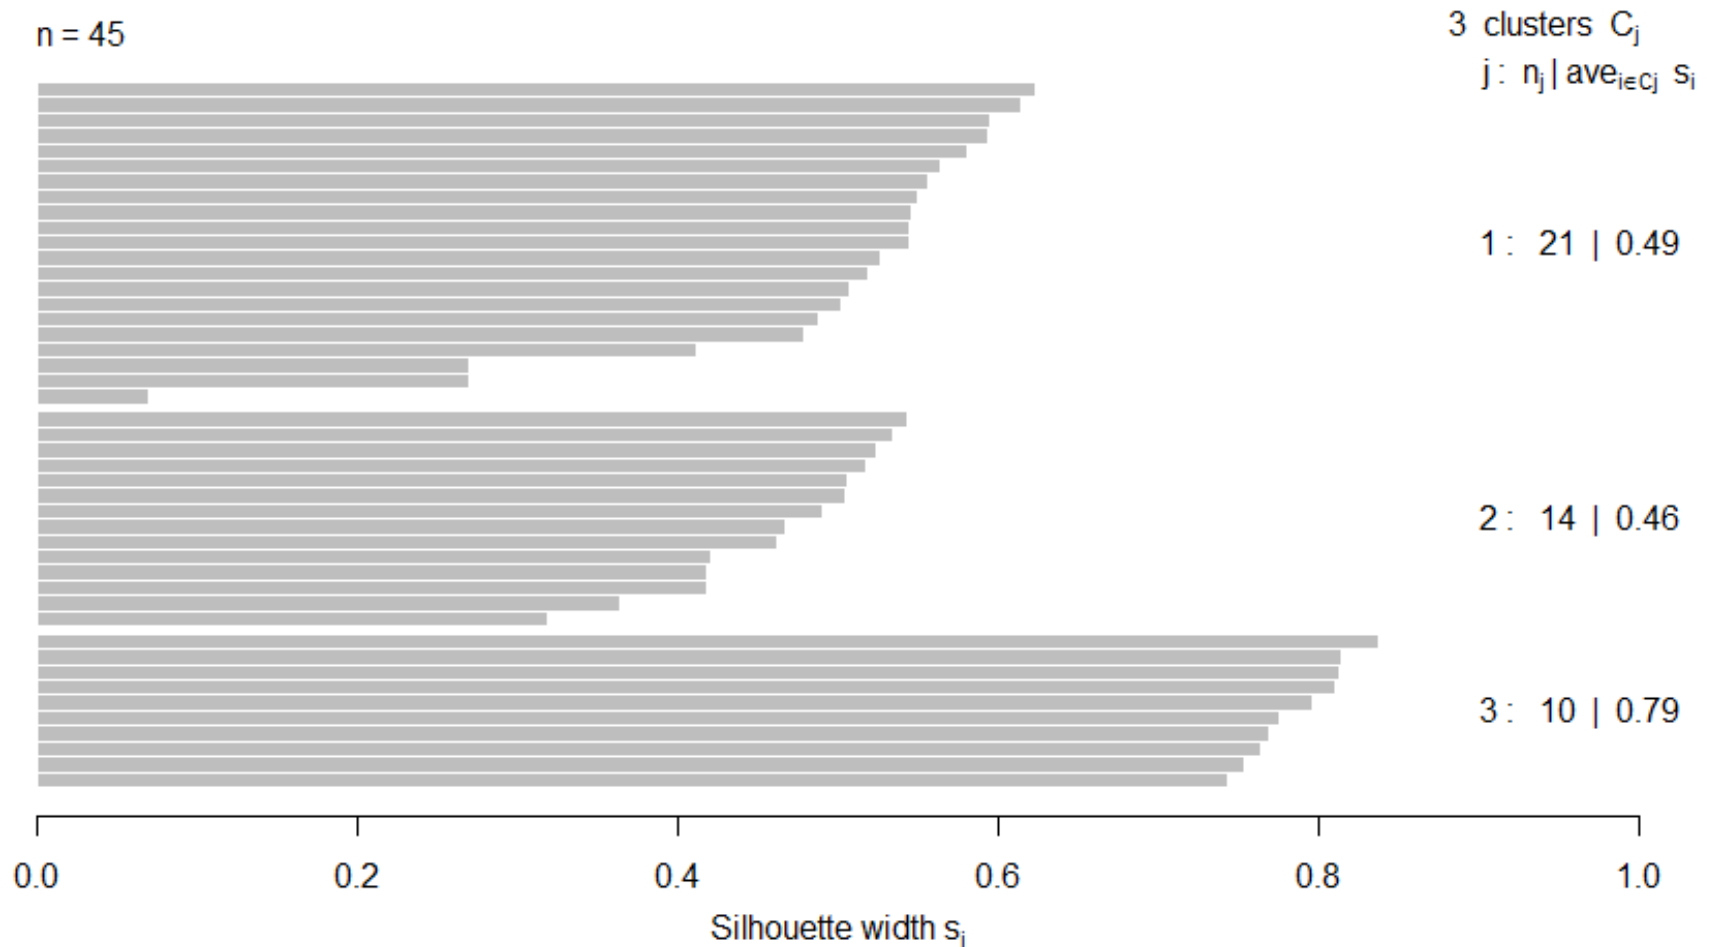
- ▶ `data(pottery)` in package HSAUR2, Chemical composition of Romano-British pottery – a data frame with 45 observations on the following 9 chemicals (Al_2O_3 , Fe_2O_3 , MgO , CaO , Na_2O , K_2O , TiO_2 , MnO , BaO , *kiln* – *site at which the pottery was found*).
- 1. Scale data – calculate Euclidean distance without *kiln* (`distpottery`)
- 2. Apply agglomerative clustering using complete linkage (`cc = hclust(distpottery, method = 'complete')`)
- 3. Split into 3 groups (`grps = cutree(cc, k = 3)`)
- 4. Plot the tree and use Silhouette plot to check the groups (`plot(silhouette(grps, distpottery))`)
- 5. Visualize results in PC1 and PC2




j: cluster; n_j : no. of members

Silhouette plot

n = 45



Average silhouette width : 0.55



agnes() in package cluster

- Can directly use the data set

```
> head(iris)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa

- Two ways to calculate distance: `dist()` and `daisy()`

```
> d = dist(iris.use)
```

```
> z = agnes(d)
```

```
> plot(z) # use which.plots= to specify which one only to display
```

```
> Hit <Return> to see next plot:
```

Example: *Iris*

data(iris)

Iris Setosa



Iris Versicolor




Iris Virginica



plot(z)


Plot(z, which.plot = 2) # only dendrogram

Banner of agnes(x = d)



Agglomerative Coefficient = 0.93

Dendrogram of agnes(x = d)



Agglomerative Coefficient = 0.93

agnes() in package cluster (cont'd)

```
> table(cutree(z,3),iris$Species)
```

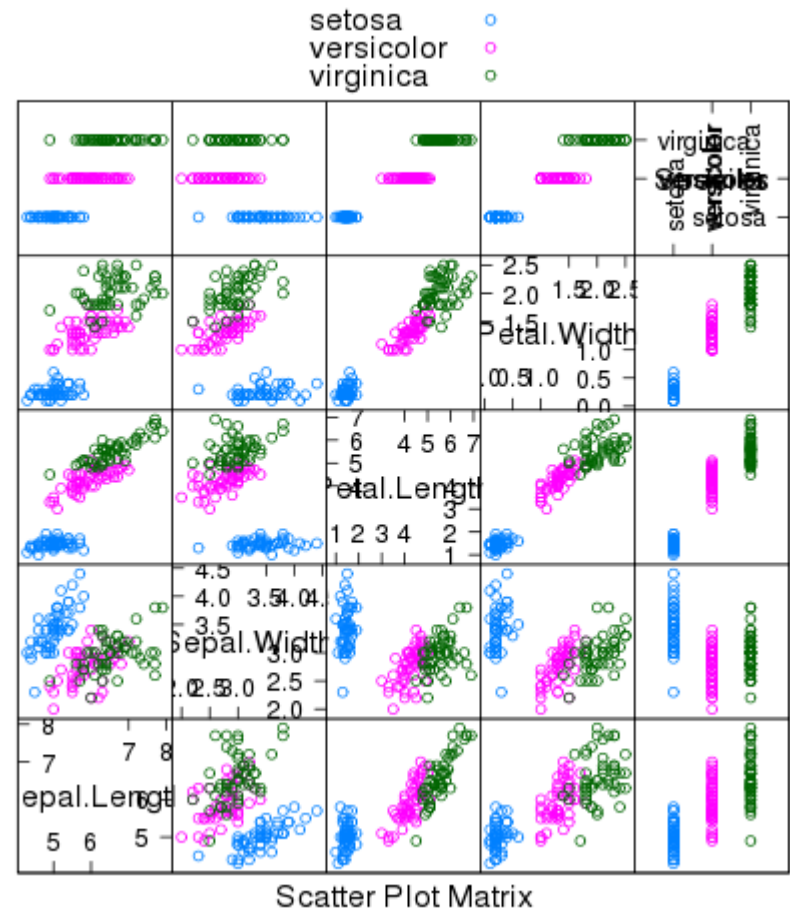
	setosa	versicolor	virginica
--	--------	------------	-----------

1	50	0	0
---	----	---	---

2	0	50	14
---	---	----	----

3	0	0	36
---	---	---	----

```
> splom(~iris, groups=iris$Species,  
auto.key=TRUE)
```



Types of Clustering of our focus



1. Hierarchical techniques

2. K-means clustering

Pre-Specify number of classes

- Number of clusters K is fixed in advance
 - Find K cluster centers μ_i and assignments, so that within-groups Sum of Squares (WGSS) is minimal
- 组内差异

$$WGSS = \sum_{all\ Cluster\ C} \sum_{Point\ i\ in\ Cluster\ C} (x_i - \mu_i)^2$$



3. Model-based clustering

K-means

- Exact solution computationally infeasible

Table 6.2: Number of possible partitions depending on the sample size n and number of clusters k .

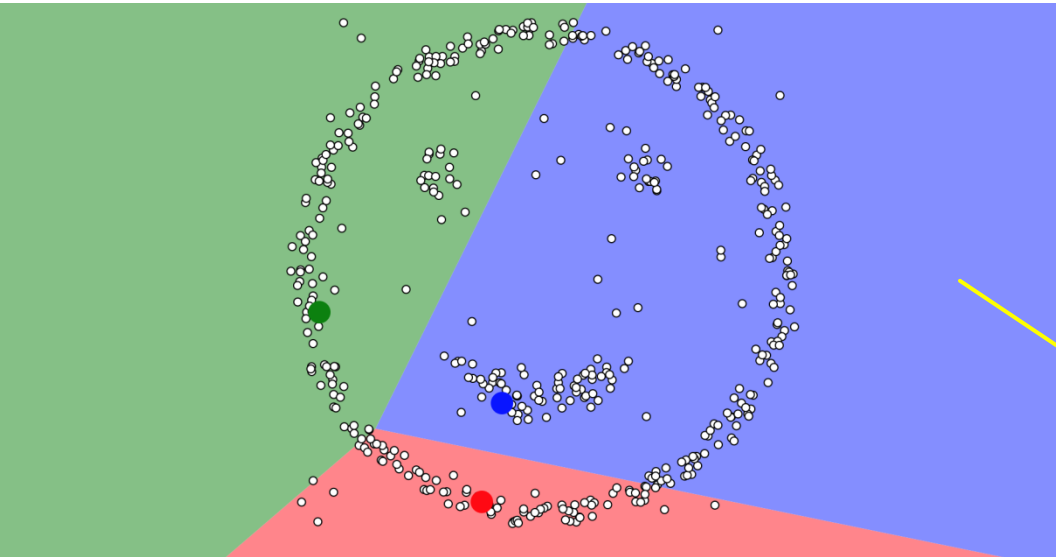
n	k	Number of possible partitions
15	3	2,375,101
20	4	45,232,115,901
25	8	690,223,721,118,368,580
100	5	10^{68}



Textbook: IAMA

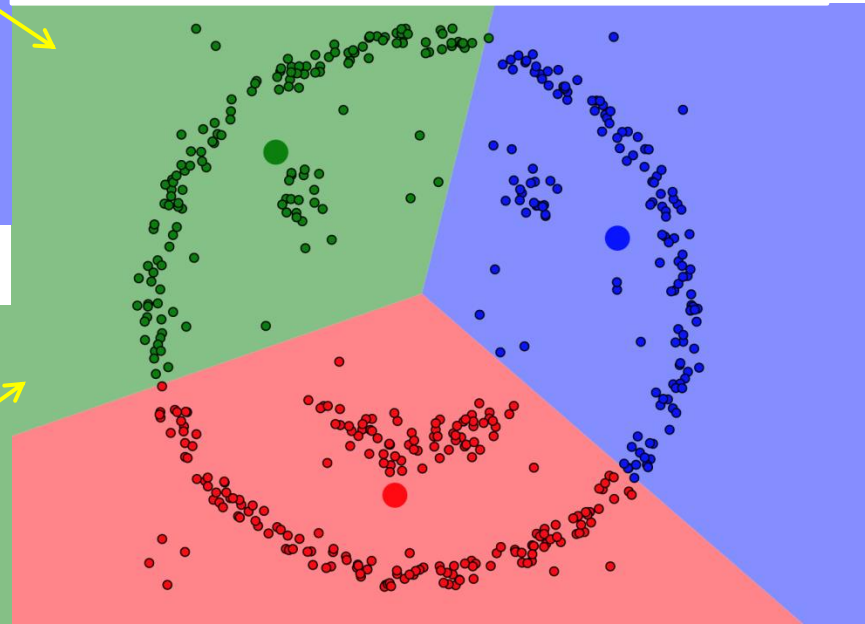
- Essential steps:
 1. Find some initial partition, with k groups
 2. Calculate the change in the clustering criterion produced by “moving” each member
 3. Keep the change that leads to the greatest improvement
 4. Repeat 2. and 3. until no useful improvement

Start #1




<http://www.naftaliharris.com/blog/visualizing-k-means-clustering/>

Clusters



Start #2



<http://stanford.edu/class/ee103/visualizations/kmeans/kmeans.html>

<http://www.naftaliharris.com/blog/visualizing-dbscan-clustering/>

kmeans() in R

- `data(pottery)` in package HSAUR2, Chemical composition of Romano-British pottery – a data frame with 45 observations on the following 9 chemicals (Al_2O_3 , Fe_2O_3 , MgO , CaO , Na_2O , K_2O , TiO_2 , MnO , BaO , kiln – *site at which the pottery was found*).

```
> ckm = kmeans(pots, centers = 3, nstart = 10)
```

```
> grpsKM = ckm$cluster
```


```
> grpsKM
```

[illegible]

kmeans() in R (cont'd)

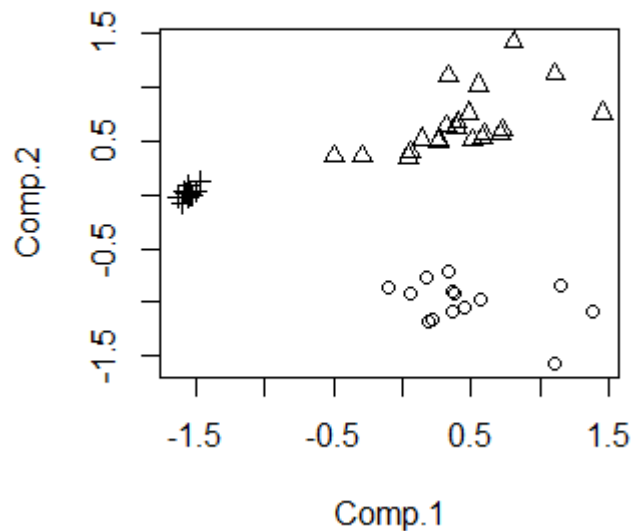
```
## find suitable number of centers
n = nrow(pottery); wss = rep(0, 6);
for (i in 1:6) wss[i] = sum(kmeans(pots, centers =
i)$withinss)
plot(1:6, wss, type = "b", xlab = "Number of groups", ylab =
"Within groups sum of squares")
```

Result may vary,
because of random
starting
configurations in
kmeans



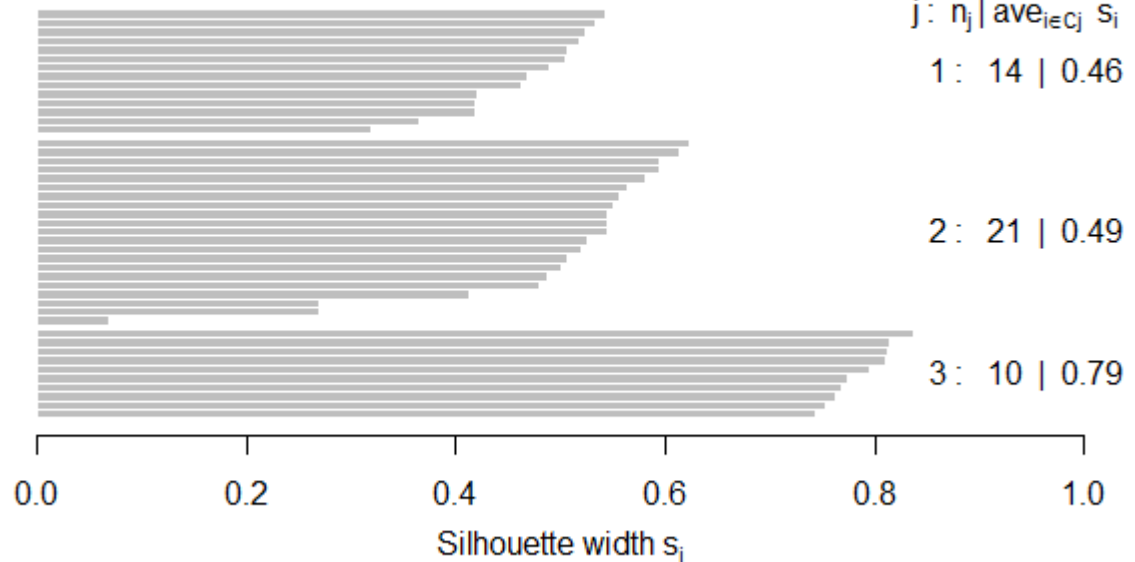
kmeans() in R (cont'd)

```
## Silhouette Plot  
plot(silhouette(grpsKM, dp))  
## visualize in PC 1 & 2  
pr = princomp(pots)$scores[,1:2]  
plot(pr, pch = grpsKM)
```



Silhouette plot of (x = grpsKM, dist = dp)

n = 45



Average silhouette width : 0.55

pam() in package cluster

► The k-Medoids Clustering

- A cluster is represented with the object closest to the center of the cluster;
- More robust in presence of outliers

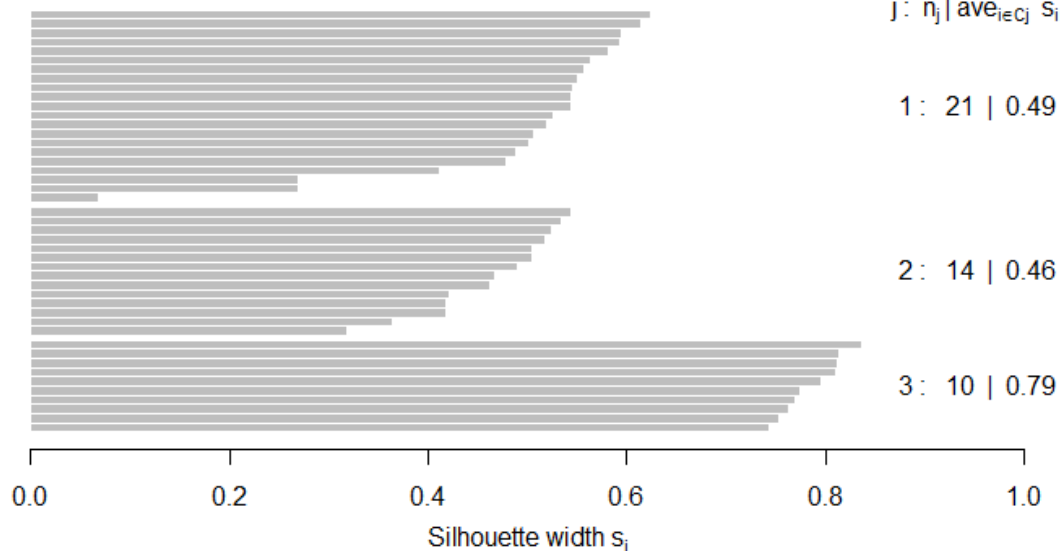
```
> pamC = pam(x = distpottery, k = 3)
> pamC$medoids
[1] "2" "26" "38"
> pamC$clustering
1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38
1 1 2 2 2 2 2 2 2 2 2 2 2 2 3 3 3
39 40 41 42 43 44 45
3 3 3 3 3 3 3
```


pam() in package cluster (cont'd)

```
> plot(pamC) ## Plots Shilouette directly
```

Silhouette plot of pam(x = dp, k = 3)

n = 45



Average silhouette width : 0.55

Types of Clustering of our focus

1. Hierarchical techniques
2. K-means clustering
3. Model-based clustering
 - Assume underlying statistical model of the population, from which we sampled our data;
 - Model assumes this population consists of a number of subpopulations (clusters)
 - Each subpopulation has variables with a different multivariate probability density function, together they result in a *finite mixture density* for the population as a whole
 - Cluster analysis → estimation of parameters of the assumed mixture
 - Determine number of clusters → model selection

Model-based clustering

- Read IAMA Ch 6.5
- R package *mclust*
 - *Model-based hierarchical clustering*
 - *Expectation-Maximization for Gaussian mixture models*
 - *Bayesian Information Criterion*

mclust

- `data(faithful)`: Waiting time between eruptions and the duration of the eruption for the Old Faithful geyser in Yellowstone National Park, Wyoming, USA.

```
> faithfulMclust <- Mclust(faithful)
> summary(faithfulMclust)
```


Gaussian finite mixture model fitted by EM algorithm

Mclust EEE (ellipsoidal, equal volume, shape and orientation) model with 3 components:

log.likelihood	n	df	BIC
-1126.4	272	11	-2314.4

Clustering table:

1	2	3
130	97	45



mclust (cont'd)

```
> summary(faithfulMclust, parameters = TRUE)
```

```
-----  
Gaussian finite mixture model fitted by EM algorithm  
-----
```

Mclust EEE (ellipsoidal, equal volume, shape and orientation) model with 3 components:

```
log.likelihood   n df      BIC  
      -1126.4 272 11 -2314.4
```

Clustering table:

```
  1  2  3  
130 97 45
```

Mixing probabilities:

```
      1      2      3  
0.46190 0.35646 0.18164
```

Means:

```
      [,1] [,2] [,3]  
eruptions 4.4761 2.0378 3.8199  
waiting   80.8922 54.4935 77.6711
```

Variances:

[:,1]

```
eruptions 0.07728 0.4765  
waiting   0.47650 33.7485
```


[:,2]

```
eruptions 0.07728 0.4765  
waiting   0.47650 33.7485
```


[:,3]

```
eruptions 0.07728 0.4765  
waiting   0.47650 33.7485
```


```
> plot(faithfulMclust)
```



Uncertainty




Classification



looks like we could fit an LDA

log Density Contour Plot



The covariance structures defining the models available in mclust

identifier	Model	HC	EM	Distribution	Volume	Shape	Orientation
E		•	•	(univariate)	equal		
V		•	•	(univariate)	variable		
EII	λI	•	•	Spherical	equal	equal	NA
VII	$\lambda_k I$	•	•	Spherical	variable	equal	NA
EEI	λA		•	Diagonal	equal	equal	coordinate axes
VEI	$\lambda_k A$		•	Diagonal	variable	equal	coordinate axes
EVI	λA_k		•	Diagonal	equal	variable	coordinate axes
VVI	$\lambda_k A_k$		•	Diagonal	variable	variable	coordinate axes
EEE	$\lambda D A D^T$	•	•	Ellipsoidal	equal	equal	equal
EEV	$\lambda D_k A D_k^T$		•	Ellipsoidal	equal	equal	variable
VEV	$\lambda_k D_k A D_k^T$		•	Ellipsoidal	variable	equal	variable
VVV	$\lambda_k D_k A_k D_k^T$	•	•	Ellipsoidal	variable	variable	variable

mclustBIC()

```
> faithfulBIC <- mclustBIC(faithful)
> faithfulSummary <- summary(faithfulBIC, data = faithful)
> faithfulSummary
```

classification table:

	1	2	3
	130	97	45

best BIC values:

	EEE,3	EEE,4	VVV,2
	-2314.4	-2320.2	-2322.2

```
> faithfulBIC
```

BIC:

	EII	VII	EEI	VEI	EVI	VVI	EEE	EEV	VEV	VVV
1	-4024.7	-4024.7	-3055.8	-3055.8	-3055.8	-3055.8	-2607.6	-2607.6	-2607.6	-2607.6
2	-3453.0	-3458.3	-2354.6	-2350.6	-2352.6	-2346.1	-2325.2	-2329.1	-2325.4	-2322.2
3	-3377.7	-3336.5	-2323.0	-2332.7	-2332.2	-2342.4	-2314.4	-2339.0	-2329.4	-2333.9
4	-3230.2	-3245.7	-2323.7	-2331.8	-2334.8	-2343.1	-2320.2	-2336.8	-2342.5	-2359.2
5	-3149.4	-3128.2	-2337.7	-2348.3	-2355.9	-2374.3	-2337.0	-2356.2	-2366.2	-2385.3
6	-3081.4	-3067.6	-2338.1	-2363.1	-2357.7	-2372.7	-2347.3	-2371.7	-2387.4	-2399.0
7	-2990.3	-2998.5	-2356.5	-2370.1	-2375.9	-2393.1	-2361.2	-2393.0	-2384.2	-2426.5
8	-2978.1	-2991.9	-2371.8	NA	-2396.0	NA	-2376.9	-2385.8	-2404.9	-2435.0
9	-2899.8	-2921.0	-2388.6	NA	-2399.1	NA	-2393.7	-2418.3	-2428.4	-2447.3

mclustBIC() (cont'd)

```
> plot(faithfulBIC, G = 1:7, ylim = c(-2500,-2300),  
      legendArgs = list(x = "bottomright", ncol = 5))
```

